



# CFDS® – Chartered Financial Data Scientist Introduction to Python

Prof. Dr. Natalie Packham

27 November and 11 December 2024

## Table of Contents

- ▼ [4 Financial Time Series](#)
  - [4.1 Financial Data](#)
  - [4.2 Correlation analysis and linear regression](#)

## 4 Financial Time Series

- Time series are ubiquitous in finance.
- `pandas` is the main library in Python to deal with time series.

### 4.1 Financial Data

#### Financial data

- For the time being we work with locally stored data files.
- These are in `.csv` -files (comma-separated values), where the data entries in each row are separated by commas.
- Some initialisation:

In [ ]:

```
import matplotlib.pyplot as plt
import seaborn as sns
```

#### Data import

- `pandas` provides a number of different functions and `DataFrame` methods for importing and exporting data.
- Here we use `pd.read_csv()`.
- The file that we load contains end-of-day data for different financial instruments retrieved from Thomson Reuters.

In [ ]:

```
# If using colab, then uncomment the line below and comment the line after that
#filename = 'https://raw.githubusercontent.com/packham/Python_CFDS/main/data/tr_eikon_eod_data.csv'
filename = './data/tr_eikon_eod_data.csv' # path and filename
f = open(filename, 'r') # this will give an error when using colab; just ignore it
f.readlines()[:5] # show first five lines
```

## Data import

In [ ]:

```
data = pd.read_csv(filename, # import csv-data into DataFrame
                   index_col=0, # take first column as index
                   parse_dates=True) # index values are datetime
```

In [ ]:

```
data.info() # information about the DataFrame object
```

## Data import

In [ ]:

```
data.head()
```

## Data import

In [ ]:

```
data.tail()
```

## Data import

In [ ]:

```
data.plot(figsize=(10, 10), subplots=True);
```

## Data import

- The identifiers used by Thomson Reuters are so-called RIC's.
- The financial instruments in the data set are:

In [ ]:

```
instruments = ['Apple Stock', 'Microsoft Stock',
               'Intel Stock', 'Amazon Stock', 'Goldman Sachs Stock',
               'SPDR S&P 500 ETF Trust', 'S&P 500 Index',
               'VIX Volatility Index', 'EUR/USD Exchange Rate',
               'Gold Price', 'VanEck Vectors Gold Miners ETF',
               'SPDR Gold Trust']
```

## Data import

In [ ]:

```
for ric, name in zip(data.columns, instruments):
    print('{:8s} | {}'.format(ric, name))
```

## Summary statistics

In [ ]:

```
data.describe().round(2)
```

## Summary statistics

- The `aggregate()` -function allows to customise the statistics viewed:

In [ ]:

```
data.aggregate([min,
                 np.mean,
                 np.std,
                 np.median,
                 max]
               ).round(2)
```

## Returns

- When working with financial data we typically (=always - you must have good reasons to deviate from this) work with performance data, i.e., **returns**.
- Reasoning:
  - Historical data are mainly used to make forecasts one or several time periods forward.
  - The daily average stock price over the last eight years is meaningless to make a forecast for tomorrow's stock price.
  - However, the daily returns are possible scenarios for the next time period(s).
- The function `pct_change()` calculates discrete returns:

$$r_t^d = \frac{S_t - S_{t-1}}{S_{t-1}},$$

where  $S_t$  denotes the stock price at time  $t$ .

## Returns

In [ ]:

```
data.pct_change().round(3).head()
```

## Returns

In [ ]:

```
data.pct_change().mean().plot(kind='bar', figsize=(10, 6));
```

## Returns

- In finance, **log-returns**, also called **continuous returns**, are often preferred over discrete returns:

$$r_t^c = \ln\left(\frac{S_t}{S_{t-1}}\right).$$

- The main reason is that log-returns are additive over time.
- For example, the log-return from  $t - 1$  to  $t + 1$  is the sum of the single-period log-returns:

$$r_{t-1,t+1}^c = \ln\left(\frac{S_{t+1}}{S_t}\right) + \ln\left(\frac{S_t}{S_{t-1}}\right) = \ln\left(\frac{S_{t+1}}{S_t} \cdot \frac{S_t}{S_{t-1}}\right) = \ln\left(\frac{S_{t+1}}{S_{t-1}}\right).$$

- Note: If the sampling (time) interval is small (e.g. one day or one week), then the difference between discrete returns and log-returns is negligible.

## Returns

In [ ]:

```
rets = np.log(data / data.shift(1)) # calculates log-returns in a vectorised way
```

In [ ]:

```
rets.head().round(3)
```

## Returns

In [ ]:

```
rets.cumsum().apply(np.exp).plot(figsize=(10, 6)); # recover price paths from log-returns
```

## 4.2 Correlation analysis and linear regression

- To further illustrate how to work with financial time series we consider the S&P 500 stock index and the VIX volatility index.
- Empirical stylised fact: As the S&P 500 rises, the VIX falls, and vice versa.
- Note: This is about **correlation** not **causation**.

## Correlation analysis

In [ ]:

```
# EOD data from Thomson Reuters Eikon Data API

# If using colab, then uncomment the line below and comment the line after that
#raw = pd.read_csv('https://raw.githubusercontent.com/packham/Python_CFDS/main/data/
raw = pd.read_csv('./data/tr_eikon_eod_data.csv', index_col=0, parse_dates=True)
data = raw[['SPX', 'VIX']].dropna()
data.tail()
```

## Correlation analysis

In [ ]:

```
data.plot(subplots=True, figsize=(10, 6));
```

## Correlation analysis

- Transform both data series into log-returns:

In [ ]:

```
rets = np.log(data / data.shift(1))
rets.head()
```

In [ ]:

```
rets.dropna(inplace=True) # drop NaN (not-a-number) entries
```

## Correlation analysis

In [ ]:

```
rets.plot(subplots=True, figsize=(10, 6));
```

## Correlation analysis

In [ ]:

```
pd.plotting.scatter_matrix(rets,
                            alpha=0.2,
                            diagonal='hist',
                            hist_kws={'bins': 35},
                            figsize=(10, 6));
```

## Correlation analysis

In [ ]:

```
rets.corr()
```

## OLS regression

- **Linear regression** captures the linear relationship between two variables.
- For two variables  $x, y$ , we postulate a linear relationship:

$$y = \alpha + \beta x + \varepsilon, \quad \alpha, \beta \in \mathbb{R}.$$

- Here,  $\alpha$  is the **intercept**,  $\beta$  is the **slope (coefficient)** and  $\varepsilon$  is the **error term**.
- Given data sample of joint observations  $(x_1, y_1), \dots, (x_n, y_n)$ , we set

$$y_i = \hat{\alpha} + \hat{\beta}x_i + \hat{\varepsilon}_i,$$

where  $\hat{\alpha}$  and  $\hat{\beta}$  are estimates of  $\alpha, \beta$  and  $\hat{\varepsilon}_1, \dots, \hat{\varepsilon}_n$  are the so-called **residuals**.

- The **ordinary least squares (OLS)** estimator  $\hat{\alpha}, \hat{\beta}$  corresponds to those values of  $\alpha, \beta$  that minimise the sum of squared residuals:

$$\min_{\alpha, \beta} \sum_{i=1}^n \varepsilon_i^2 = \sum_{i=1}^n (y_i - \alpha - \beta x_i)^2.$$

## OLS regressions

- Simplest form of OLS regression:

In [ ]:

```
reg = np.polyfit(rets['.SPX'], rets['.VIX'], deg=1) # fit a linear equation (a poly
reg.view() # the fitted paramters
```

2.62e-03 is scientific notation:  $2.62e - 03 = 2.62 \cdot 10^{-3}$ .

In [ ]:

```
ax = rets.plot(kind='scatter', x='.SPX', y='.VIX', figsize=(8, 5))
ax.plot(rets['.SPX'], np.polyval(reg, rets['.SPX']), 'r', lw=2);
```

## OLS regression

- To do a more refined OLS regression with a proper analysis, use the package `statsmodels`.

In [ ]:

```
import statsmodels.api as sm

Y=rets['.VIX']
X=rets['.SPX']
X = sm.add_constant(X)
```

In [ ]:

```
model = sm.OLS(Y,X)
results = model.fit()
```

In [ ]:

```
results.params
```

In [ ]:

```
results.predict()[0:10]
```

## OLS regression

In [ ]:

```
print(results.summary())
```

### OLS regression: Interpretation of output and forecasting

- The column `coef` lists the coefficients of the regression: the coefficient in the row labelled `const` corresponds to  $\hat{\alpha}$  ( $= 0.0026$ ) and the coefficient in the row `.SPX` denotes  $\hat{\beta}$  ( $= -6.6515$ ).
- The estimated model in the example is thus:  

$$.VIX = 0.0026 - 6.6516.SPX.$$
- The best forecast of the VIX return when observing an S&P return of 2% is therefore  $0.0026 - 6.6516 \cdot 0.02 = -0.130432 = -13.0432\%$ .

### OLS regression: Validation ( $R^2$ )


- To **validate** the model, i.e., to determine, if the model in itself and the explanatory variable(s) make sense, we look  $R^2$  and various  $p$ -values (or confidence intervals or  $t$ -statistics).
- $R^2$  measures the fraction of variance in the dependent variable  $Y$  that is captured by the regression line;  $1 - R^2$  is the fraction of  $Y$ -variance that remains in the residuals  $\varepsilon_i^2$ ,  $i = 1, \dots, n$
- In the output above  $R^2$  is given as 0.647. In other words, 64.7% of the variance in VIX returns are "explained" by SPX returns.
- A high  $R^2$  (and this one is high) is necessary for making forecasts.

### OLS regression: Validation (confidence interval)

- An important hypothesis to test in any regression model is whether the explanatory variable(s) have an effect on the independent variable.
- This can be translated into testing whether  $\beta \neq 0$ . ( $\beta = 0$  is the same as saying that the  $X$  variable can be removed from the model.)
- Formally, we test the null hypothesis  $H_0 : \beta = 0$  against the alternative hypothesis  $H_1 : \beta \neq 0$ .
- There are several statistics to come to the same conclusion: confidence intervals,  $t$ -statistics and  $p$ -values.
- The **confidence interval** is an interval around the estimate  $\hat{\beta}$  that we are confident contains the true parameter  $\beta$ . A typical **confidence level** is 95%.
- If the 95% confidence interval does **not** contain 0, then we say  $\beta$  is **statistically significant** at the 5% ( $=1-95\%$ ) level, and we conclude that  $\beta \neq 0$ .

### OLS regression: Validation ( $t$ -statistic)

- The  $t$ -statistic corresponds to the **number of standard deviations** that the estimated coefficient  $\hat{\beta}$  is away from 0 (the mean under  $H_0$ ).
- For a normal distribution, we have the following rules of thumb:
  - 66% of observations lie within one standard deviation of the mean
  - 95% of observations lie within two standard deviations of the mean
  - 99.7% of observations lie within three standard deviations of the mean



- If the sample size is large enough, then the  $t$ -statistic is approximately normally distributed, and if it is large (in absolute terms), then this is an indication against  $\beta = 0$ .
- In the example above, the  $t$ -statistics is -62.559, i.e.,  $\hat{\beta}$  is approx. 63 standard deviations away from zero, which is practically impossible.

## OLS regression: Validation ( $p$ -value)

- The  $p$ -value expresses the probability of observing a coefficient estimate as extreme (away from zero) as  $\hat{\beta}$  under  $H_0$ , i.e., when  $\beta = 0$ .
- In other words, it measures the probability of observing a  $t$ -statistic as extreme as the one observed if  $\beta = 0$ .
- If the  $p$ -value (column  $P > |t|$ ) is smaller than the desired level of significance (typically 5%), then the  $H_0$  can be rejected and we conclude that  $\beta \neq 0$ .
- In the example above, the  $p$ -value is given as 0.000, i.e., it is so small, that we can conclude the estimated coefficient  $\hat{\beta}$  is so extreme (= away from zero) that is virtually impossible to obtain such an estimated if  $\beta = 0$ .
- Finally, the  $F$ -test tests the hypotheses  $H_0 : R^2 = 0$  versus  $H_1 : R^2 \neq 0$ . In a multiple regression with  $k$  independent variables, this is equivalent to  $H_0 : \beta_1 = \dots = \beta_k = 0$ .
- In the example above, the  $p$ -value of the  $F$ -test is 0, so we conclude that the model overall has explanatory power.