



**CFDS® – Chartered Financial Data Scientist**

# **Introduction to Python**

**Prof. Dr. Natalie Packham**

**29 November and 13 December 2023**

# Table of Contents

- 4 Financial Time Series
  - 4.1 Financial Data
  - 4.2 Correlation analysis and linear regression

# Financial Time Series

- Time series are ubiquitous in finance.
- `pandas` is the main library in Python to deal with time series.

Financial Data

## Financial data

- For the time being we work with locally stored data files.
- These are in `.csv`-files (comma-separated values), where the data entries in each row are separated by commas.
- Some initialisation:

```
In [31]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
plt.style.use('seaborn')
plt.rcParams['font.family'] = 'serif'
```

```
/var/folders/46/b127yp714m71zfmt9j7_lhwh0000gq/T/ipykernel_5169
8/2492143664.py:4: MatplotlibDeprecationWarning: The seaborn st
yles shipped by Matplotlib are deprecated since 3.6, as they no
longer correspond to the styles shipped by seaborn. However, th
ey will remain available as 'seaborn-v0_8-<style>'. Alternative
ly, directly use the seaborn API instead.
plt.style.use('seaborn')
```

# Data import

- `pandas` provides a number of different functions and `DataFrame` methods for importing and exporting data.
- Here we use `pd.read_csv()`.
- The file that we load contains end-of-day data for different financial instruments retrieved from Thomson Reuters.

```
In [32]: # If using colab, then uncomment the line below and comment the line at
#filename = 'https://raw.githubusercontent.com/packham/Python_CFDS/main
filename = './data/tr_eikon_eod_data.csv' # path and filename
f = open(filename, 'r') # this will give an error when using colab; jus
f.readlines()[:5] # show first five lines
```

```
Out[32]: ['Date,AAPL.O,MSFT.O,INTC.O,AMZN.O,GS.N,SPY,.SPX,.VIX,EUR=XAU
=,GDX,GLD\n',
'2010-01-01,,,,,,,,,1.4323,1096.35,,\n',
'2010-01-04,30.57282657,30.95,20.88,133.9,173.08,113.33,1132.9
9,20.04,1.4411,1120.0,47.71,109.8\n',
'2010-01-05,30.625683660000004,30.96,20.87,134.69,176.14,113.6
3,1136.52,19.35,1.4368,1118.65,48.17,109.7\n',
'2010-01-06,30.138541290000003,30.77,20.8,132.25,174.26,113.7
1,1137.14,19.16,1.4412,1138.5,49.34,111.51\n']
```

# Data import

```
In [33]: data = pd.read_csv(filename, # import csv-data into DataFrame
                             index_col=0, # take first column as index
                             parse_dates=True) # index values are datetime
```

```
In [34]: data.info() # information about the DataFrame object
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 2216 entries, 2010-01-01 to 2018-06-29
Data columns (total 12 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0   AAPL.O      2138 non-null   float64
 1   MSFT.O      2138 non-null   float64
 2   INTC.O      2138 non-null   float64
 3   AMZN.O      2138 non-null   float64
 4   GS.N        2138 non-null   float64
 5   SPY         2138 non-null   float64
 6   .SPX        2138 non-null   float64
 7   .VIX        2138 non-null   float64
 8   EUR=        2216 non-null   float64
 9   XAU=        2211 non-null   float64
10   GDX         2138 non-null   float64
11   GLD         2138 non-null   float64
dtypes: float64(12)
memory usage: 225.1 KB
```



## Data import

In [35]: `data.head()`

Out[35]:

	AAPL.O	MSFT.O	INTC.O	AMZN.O	GS.N	SPY	.SPX	.VIX	EU
Date									
2010-01-01	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	1.43
2010-01-04	30.572827	30.950	20.88	133.90	173.08	113.33	1132.99	20.04	1.44
2010-01-05	30.625684	30.960	20.87	134.69	176.14	113.63	1136.52	19.35	1.43
2010-01-06	30.138541	30.770	20.80	132.25	174.26	113.71	1137.14	19.16	1.44
2010-01-07	30.082827	30.452	20.60	130.00	177.67	114.19	1141.69	19.06	1.43

# Data import


In [36]: `data.tail()`


Out[36]:

	AAPL.O	MSFT.O	INTC.O	AMZN.O	GS.N	SPY	.SPX	.VIX	EUR:
<b>Date</b>									
<b>2018-06-25</b>	182.17	98.39	50.71	1663.15	221.54	271.00	2717.07	17.33	1.170%
<b>2018-06-26</b>	184.43	99.08	49.67	1691.09	221.58	271.60	2723.06	15.92	1.164%
<b>2018-06-27</b>	184.16	97.54	48.76	1660.51	220.18	269.35	2699.63	17.91	1.155%
<b>2018-06-28</b>	185.50	98.63	49.25	1701.45	223.42	270.89	2716.31	16.85	1.156%
<b>2018-06-29</b>	185.11	98.61	49.71	1699.80	220.57	271.28	2718.37	16.09	1.168%

# Data import

```
In [37]: data.plot(figsize=(10, 10), subplots=True);
```





## Data import

- The identifiers used by Thomson Reuters are so-called RIC's.
- The financial instruments in the data set are:

```
In [38]: instruments = ['Apple Stock', 'Microsoft Stock',  
                        'Intel Stock', 'Amazon Stock', 'Goldman Sachs Stock',  
                        'SPDR S&P 500 ETF Trust', 'S&P 500 Index',  
                        'VIX Volatility Index', 'EUR/USD Exchange Rate',  
                        'Gold Price', 'VanEck Vectors Gold Miners ETF',  
                        'SPDR Gold Trust']
```

## Data import

```
In [39]: for ric, name in zip(data.columns, instruments):  
         print('{:8s} | {}'.format(ric, name))
```

```
AAPL.O    | Apple Stock  
MSFT.O    | Microsoft Stock  
INTC.O    | Intel Stock  
AMZN.O    | Amazon Stock  
GS.N      | Goldman Sachs Stock  
SPY       | SPDR S&P 500 ETF Trust  
.SPX      | S&P 500 Index  
.VIX      | VIX Volatility Index  
EUR=      | EUR/USD Exchange Rate  
XAU=      | Gold Price  
GDX       | VanEck Vectors Gold Miners ETF  
GLD       | SPDR Gold Trust
```

## Summary statistics

In [40]: `data.describe().round(2)`

Out[40]:

	AAPL.O	MSFT.O	INTC.O	AMZN.O	GS.N	SPY	.SPX	.VIX
<b>count</b>	2138.00	2138.00	2138.00	2138.00	2138.00	2138.00	2138.00	2138.00
<b>mean</b>	93.46	44.56	29.36	480.46	170.22	180.32	1802.71	17.03
<b>std</b>	40.55	19.53	8.17	372.31	42.48	48.19	483.34	5.88
<b>min</b>	27.44	23.01	17.66	108.61	87.70	102.20	1022.58	9.14
<b>25%</b>	60.29	28.57	22.51	213.60	146.61	133.99	1338.57	13.07
<b>50%</b>	90.55	39.66	27.33	322.06	164.43	186.32	1863.08	15.58
<b>75%</b>	117.24	54.37	34.71	698.85	192.13	210.99	2108.94	19.07
<b>max</b>	193.98	102.49	57.08	1750.08	273.38	286.58	2872.87	48.00

## Summary statistics

- The `aggregate()` -function allows to customise the statistics viewed:

```
In [41]: data.aggregate([min,  
                        np.mean,  
                        np.std,  
                        np.median,  
                        max]  
        ).round(2)
```

```
Out[41]:
```

	AAPL.O	MSFT.O	INTC.O	AMZN.O	GS.N	SPY	.SPX	.VIX	EUF
<b>min</b>	27.44	23.01	17.66	108.61	87.70	102.20	1022.58	9.14	1.0
<b>mean</b>	93.46	44.56	29.36	480.46	170.22	180.32	1802.71	17.03	1.1
<b>std</b>	40.55	19.53	8.17	372.31	42.48	48.19	483.34	5.88	0.1
<b>median</b>	90.55	39.66	27.33	322.06	164.43	186.32	1863.08	15.58	1.1
<b>max</b>	193.98	102.49	57.08	1750.08	273.38	286.58	2872.87	48.00	1.4



# Returns

- When working with financial data we typically (=always - you must have good reasons to deviate from this) work with performance data, i.e., **returns**.
- Reasoning:
  - Historical data are mainly used to make forecasts one or several time periods forward.
  - The daily average stock price over the last eight years is meaningless to make a forecast for tomorrow's stock price.
  - However, the daily returns are possible scenarios for the next time period(s).
- The function `pct_change()` calculates discrete returns:

$$r_t^d = \frac{S_t - S_{t-1}}{S_{t-1}},$$

where  $S_t$  denotes the stock price at time  $t$ .


## Returns

```
In [42]: data.pct_change().round(3).head()
```

	AAPL.O	MSFT.O	INTC.O	AMZN.O	GS.N	SPY	.SPX	.VIX	EUR=
Date									
2010-01-01	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
2010-01-04	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	0.006
2010-01-05	0.002	0.000	-0.000	0.006	0.018	0.003	0.003	-0.034	-0.003
2010-01-06	-0.016	-0.006	-0.003	-0.018	-0.011	0.001	0.001	-0.010	0.003
2010-01-07	-0.002	-0.010	-0.010	-0.017	0.020	0.004	0.004	-0.005	-0.007

# Returns

```
In [43]: data.pct_change().mean().plot(kind='bar', figsize=(10, 6));
```



# Returns

- In finance, **log-returns**, also called **continuous returns**, are often preferred over discrete returns:  $r_t^c = \ln\left(\frac{S_t}{S_{t-1}}\right)$ .
- The main reason is that log-return are additive over time.
- For example, the log-return from  $t - 1$  to  $t + 1$  is the sum of the single-period log-returns:

$$r_{t-1,t+1}^c = \ln\left(\frac{S_{t+1}}{S_t}\right) + \ln\left(\frac{S_t}{S_{t-1}}\right) = \ln\left(\frac{S_{t+1}}{S_t} \cdot \frac{S_t}{S_{t-1}}\right) = \ln\left(\frac{S_{t+1}}{S_{t-1}}\right).$$

- Note: If the sampling (time) interval is small (e.g. one day or one week), then the difference between discrete returns and log-returns is negligible.

# Returns

```
In [44]: rets = np.log(data / data.shift(1)) # calculates log-returns in a vect
```


```
In [45]: rets.head().round(3)
```

```
Out[45]:
```

	AAPL.O	MSFT.O	INTC.O	AMZN.O	GS.N	SPY	.SPX	.VIX	EUR=
Date									
2010-01-01	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
2010-01-04	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	0.006
2010-01-05	0.002	0.000	-0.000	0.006	0.018	0.003	0.003	-0.035	-0.003
2010-01-06	-0.016	-0.006	-0.003	-0.018	-0.011	0.001	0.001	-0.010	0.003
2010-01-07	-0.002	-0.010	-0.010	-0.017	0.019	0.004	0.004	-0.005	-0.007

# Returns

```
In [46]: rets.cumsum().apply(np.exp).plot(figsize=(10, 6)); # recover price pat
```



# Correlation analysis and linear regression

- To further illustrate how to work with financial time series we consider the S&P 500 stock index and the VIX volatility index.
- Empirical stylised fact: As the S&P 500 rises, the VIX falls, and vice versa.
- Note: This is about **correlation** not **causation**.

## Correlation analysis

```
In [47]: # EOD data from Thomson Reuters Eikon Data API

# If using colab, then uncomment the line below and comment the line at
#raw = pd.read_csv('https://raw.githubusercontent.com/packham/Python_CE
raw = pd.read_csv('./data/tr_eikon_eod_data.csv', index_col=0, parse_da
data = raw[['.SPX', '.VIX']].dropna()
data.tail()
```


```
Out[47]:
```

	.SPX	.VIX
Date		
2018-06-25	2717.07	17.33
2018-06-26	2723.06	15.92
2018-06-27	2699.63	17.91
2018-06-28	2716.31	16.85
2018-06-29	2718.37	16.09



# Correlation analysis

```
In [48]: data.plot(subplots=True, figsize=(10, 6));
```



## Correlation analysis

- Transform both data series into log-returns:

```
In [49]: rets = np.log(data / data.shift(1))  
rets.head()
```

```
Out[49]:
```

	<b>.SPX</b>	<b>.VIX</b>
<b>Date</b>		
<b>2010-01-04</b>	NaN	NaN
<b>2010-01-05</b>	0.003111	-0.035038
<b>2010-01-06</b>	0.000545	-0.009868
<b>2010-01-07</b>	0.003993	-0.005233
<b>2010-01-08</b>	0.002878	-0.050024

```
In [50]: rets.dropna(inplace=True) # drop NaN (not-a-number) entries
```


# Correlation analysis

```
In [51]: rets.plot(subplots=True, figsize=(10, 6));
```



# Correlation analysis

```
In [52]: pd.plotting.scatter_matrix(rets,  
                                     alpha=0.2,  
                                     diagonal='hist',  
                                     hist_kwds={'bins': 35},  
                                     figsize=(10, 6));
```





## Correlation analysis

In [53]: `rets.corr()`

Out [53]:

	<b>.SPX</b>	<b>.VIX</b>
<b>.SPX</b>	1.000000	-0.804382
<b>.VIX</b>	-0.804382	1.000000

# OLS regression

- **Linear regression** captures the linear relationship between two variables.
- For two variables  $x, y$ , we postulate a linear relationship:

$$y = \alpha + \beta x + \varepsilon, \quad \alpha, \beta \in \mathbb{R}.$$

- Here,  $\alpha$  is the **intercept**,  $\beta$  is the **slope (coefficient)** and  $\varepsilon$  is the **error term**.
- Given data sample of joint observations  $(x_1, y_1), \dots, (x_n, y_n)$ , we set

$$y_i = \hat{\alpha} + \hat{\beta} x_i + \hat{\varepsilon}_i,$$

where  $\hat{\alpha}$  and  $\hat{\beta}$  are estimates of  $\alpha, \beta$  and  $\hat{\varepsilon}_1, \dots, \hat{\varepsilon}_n$  are the so-called **residuals**.

- The **ordinary least squares (OLS)** estimator  $\hat{\alpha}, \hat{\beta}$  corresponds to those values of  $\alpha, \beta$  that minimise the sum of squared residuals:

$$\min_{\alpha, \beta} \sum_{i=1}^n \varepsilon_i^2 = \sum_{i=1}^n (y_i - \alpha - \beta x_i)^2.$$

# OLS regressions

- Simplest form of OLS regression:


```
In [54]: reg = np.polyfit(rets['.SPX'], rets['.VIX'], deg=1) # fit a linear equation
reg.view() # the fitted parameters
```

```
Out[54]: array([-6.65160028e+00,  2.62132142e-03])
```

2.62e-03 is scientific notation:  $2.62e - 03 = 2.62 \cdot 10^{-3}$ .

```
In [55]: ax = rets.plot(kind='scatter', x='.SPX', y='.VIX', figsize=(8, 5))
ax.plot(rets['.SPX'], np.polyval(reg, rets['.SPX']), 'r', lw=2);
```





# OLS regression

- To do a more refined OLS regression with a proper analysis, use the package `statsmodels`.

```
In [56]: import statsmodels.api as sm

Y=rets['.VIX']
X=rets['.SPX']
X = sm.add_constant(X)
```

```
In [57]: model = sm.OLS(Y,X)
results = model.fit()
```

```
In [58]: results.params
```

```
Out[58]: const      0.002621
         .SPX      -6.651600
         dtype: float64
```

```
In [59]: results.predict()[0:10]
```

```
Out[59]: array([-0.01807052, -0.0010063 , -0.0239404 , -0.01651898, -0.0
0898726,
               0.06531557, -0.05252965, -0.01349928,  0.07500527, -0.0
8000615])
```

# OLS regression

```
In [60]: print(results.summary())
```

```

                                OLS Regression Results
=====
Dep. Variable:                  .VIX    R-squared:
0.647
Model:                        OLS      Adj. R-squared:
0.647
Method:                     Least Squares    F-statistic:
3914.
Date:                Sun, 26 Nov 2023    Prob (F-statistic):
0.00
Time:                15:04:17    Log-Likelihood:
3550.1
No. Observations:                2137    AIC:
-7096.
Df Residuals:                2135    BIC:
-7085.
Df Model:                        1
Covariance Type:                nonrobust
=====
=====
                                coef      std err          t      P>|t|      [0.
025      0.975]
-----
const                0.0026      0.001      2.633      0.009      0.
```

```

001          0.005
.SPX          -6.6516      0.106      -62.559      0.000      -6.
860          -6.443
=====
=====
Omnibus:                    518.582      Durbin-Watson:
2.094
Prob(Omnibus):              0.000      Jarque-Bera (JB):
6789.425
Skew:                      0.766      Prob(JB):
0.00
Kurtosis:                  11.597      Cond. No.
107.
=====
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

## OLS regression: Interpretation of output and forecasting

- The column `coef` lists the coefficients of the regression: the coefficient in the row labelled `const` corresponds to  $\hat{\alpha}$  ( $= 0.0026$ ) and the coefficient in the row `.SPX` denotes  $\hat{\beta}$  ( $= -6.6515$ ).
- The estimated model in the example is thus:

$$.VIX = 0.0026 - 6.6516.SPX.$$

- The best forecast of the VIX return when observing an S&P return of 2% is therefore  $0.0026 - 6.6516 \cdot 0.02 = -0.130432 = -13.0432\%$ .

## OLS regression: Validation ( $R^2$ )

- To **validate** the model, i.e., to determine, if the model in itself and the explanatory variable(s) make sense, we look  $R^2$  and various  $p$ -values (or confidence intervals or  $t$ -statistics).
- $R^2$  measures the fraction of variance in the dependent variable  $Y$  that is captured by the regression line;  $1 - R^2$  is the fraction of  $Y$ -variance that remains in the residuals  $\varepsilon_i^2$ ,  $i = 1, \dots, n$ .
- In the output above  $R^2$  is given as 0.647. In other words, 64.7% of the variance in VIX returns are "explained" by SPX returns.
- A high  $R^2$  (and this one is high) is necessary for making forecasts.


## OLS regression: Validation (confidence interval)

- An important hypothesis to test in any regression model is whether the explanatory variable(s) have an effect on the independent variable.
- This can be translated into testing whether  $\beta \neq 0$ . ( $\beta = 0$  is the same as saying that the  $X$  variable can be removed from the model.)
- Formally, we test the null hypothesis  $H_0 : \beta = 0$  against the alternative hypothesis  $H_1 : \beta \neq 0$ .
- There are several statistics to come to the same conclusion: confidence intervals,  $t$ -statistics and  $p$ -values.
- The **confidence interval** is an interval around the estimate  $\hat{\beta}$  that we are confident contains the true parameter  $\beta$ . A typical **confidence level** is 95%.
- If the 95% confidence interval does **not** contain 0, then we say  $\beta$  is **statistically significant** at the 5% (=1-95%) level, and we conclude that  $\beta \neq 0$ .

OLS regression: Validation ( $t$ -statistic)



- The  $t$ -statistic corresponds to the **number of standard deviations** that the estimated coefficient  $\hat{\beta}$  is away from 0 (the mean under  $H_0$ ).
- For a normal distribution, we have the following rules of thumb:
  - 66% of observations lie within one standard deviation of the mean
  - 95% of observations lie within two standard deviations of the mean
  - 99.7% of observations lie within three standard deviations of the mean



- If the sample size is large enough, then the  $t$ -statistic is approximately normally distributed, and if it is large (in absolute terms), then this is an indication against  $\beta = 0$ .
- In the example above, the  $t$ -statistics is -62.559, i.e.,  $\hat{\beta}$  is approx. 63 standard deviations away from zero, which is practically impossible.

## OLS regression: Validation ( $p$ -value)

- The  $p$ -value expresses the probability of observing a coefficient estimate as extreme (away from zero) as  $\hat{\beta}$  under  $H_0$ , i.e., when  $\beta = 0$ .
- In other words, it measures the probability of observing a  $t$ -statistic as extreme as the one observed if  $\beta = 0$ .
- If the  $p$ -value (column `P>|t|`) is smaller than the desired level of significance (typically 5%), then the  $H_0$  can be rejected and we conclude that  $\beta \neq 0$ .
- In the example above, the  $p$ -value is given as 0.000, i.e., it is so small, that we can conclude the estimated coefficient  $\hat{\beta}$  is so extreme (= away from zero) that is virtually impossible to obtain such an estimated if  $\beta = 0$ .
- Finally, the  $F$ -test tests the hypotheses  $H_0 : R^2 = 0$  versus  $H_1 : R^2 \neq 0$ . In a multiple regression with  $k$  independent variables, this is equivalent to  $H_0 : \beta_1 = \dots = \beta_k = 0$ .
- In the example above, the  $p$ -value of the  $F$ -test is 0, so we conclude that the model overall has explanatory power.